

Healthcare-insurance-analysis-project-1

```
[81]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[82]: import warnings
warnings.filterwarnings("ignore")
```

```
[83]: hospital = pd.read_csv("Hospitalisation details.csv")
medical = pd.read_csv("Medical Examinations.csv")
customer = pd.read_excel("Names.xlsx")
```

```
[84]: hospital.head()
```

```
[84]:   Customer ID  year month  date  children  charges Hospital tier City tier \
0      Id2335  1992   Jul    9         0   563.84      tier - 2 tier - 3 \
1      Id2334  1992  Nov   30         0   570.62      tier - 2 tier - 1
2      Id2333  1993   Jun   30         0   600.00      tier - 2 tier - 1
3      Id2332  1992   Sep   13         0   604.54      tier - 3 tier - 3
4      Id2331  1998   Jul   27         0   637.26      tier - 3 tier - 3

      State ID
0      R1013
1      R1013
2      R1013
3      R1013
4      R1013
```

```
[85]: hospital.shape
```

```
[85]: (2343, 9)
```

```
[86]: hospital.describe()
```

```
[86]:
```

	date	children	charges
count	2343.000000	2343.000000	2343.000000
mean	15.554844	1.026035	13559.067870

std	8.721194	1.233847	11922.658415
min	1.000000	0.000000	563.840000
25%	8.000000	0.000000	5084.010000
50%	15.000000	0.000000	9634.540000
75%	23.000000	2.000000	17029.675000
max	30.000000	5.000000	63770.430000

```
[87]: hospital.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2343 entries, 0 to 2342
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Customer ID     2343 non-null   object
1   year            2343 non-null   object
2   month           2343 non-null   object
3   date            2343 non-null   int64
4   children        2343 non-null   int64
5   charges         2343 non-null   float64
6   Hospital tier   2343 non-null   object
7   City tier        2343 non-null   object
8   State ID        2343 non-null   object
dtypes: float64(1), int64(2), object(6)
memory usage: 164.9+ KB
```

```
[88]: medical.head()
```

```
[88]:   Customer ID    BMI  HBA1C Heart Issues Any Transplants Cancer history
0         Id1  47.410   7.47           No           No           No \
1         Id2  30.360   5.77           No           No           No
2         Id3  34.485  11.87          yes           No           No
3         Id4  38.095   6.05           No           No           No
4         Id5  35.530   5.45           No           No           No

   NumberOfMajorSurgeries  smoker
0      No major surgery     yes
1      No major surgery     yes
2                2         yes
3      No major surgery     yes
4      No major surgery     yes
```

```
[89]: medical.shape
```

```
[89]: (2335, 8)
```

```
[90]: medical.describe()
```

```
[90]:
```

	BMI	HBA1C
count	2335.000000	2335.000000
mean	30.972649	6.578998
std	8.742095	2.228731
min	15.010000	4.000000
25%	24.600000	4.900000
50%	30.400000	5.810000
75%	36.300000	7.955000
max	55.050000	12.000000

```
[91]: medical.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2335 entries, 0 to 2334
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Customer ID                          2335 non-null   object
1   BMI                                  2335 non-null   float64
2   HBA1C                               2335 non-null   float64
3   Heart Issues                        2335 non-null   object
4   Any Transplants                     2335 non-null   object
5   Cancer history                      2335 non-null   object
6   NumberOfMajorSurgeries              2335 non-null   object
7   smoker                             2335 non-null   object
dtypes: float64(2), object(6)
memory usage: 146.1+ KB
```

```
[92]: customer.head()
```

```
[92]:
```

	Customer ID	name
0	Id1	Hawks, Ms. Kelly
1	Id2	Lehner, Mr. Matthew D
2	Id3	Lu, Mr. Phil
3	Id4	Osborne, Ms. Kelsey
4	Id5	Kadala, Ms. Kristyn

```
[93]: customer.shape
```

```
[93]: (2335, 2)
```

```
[94]: customer.describe()
```

```
[94]:
```

	Customer ID	name
count	2335	2335
unique	2335	2335
top	Id1	Hawks, Ms. Kelly

```
freq          1          1
```

```
[95]: customer.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2335 entries, 0 to 2334
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Customer ID 2335 non-null   object
1   name        2335 non-null   object
dtypes: object(2)
memory usage: 36.6+ KB
```

0.1 1. Collate the files so that all the information is in one place

```
[96]: df = pd.merge(pd.merge(hospital,medical,on='Customer ID'),customer,on='Customer_ID')
```

```
[97]: df
```

```
[97]:
```

	Customer ID	year	month	date	children	charges	Hospital	tier
0	Id2335	1992	Jul	9	0	563.84	tier - 2	\
1	Id2334	1992	Nov	30	0	570.62	tier - 2	
2	Id2333	1993	Jun	30	0	600.00	tier - 2	
3	Id2332	1992	Sep	13	0	604.54	tier - 3	
4	Id2331	1998	Jul	27	0	637.26	tier - 3	
...
2330	Id5	1989	Jun	19	0	55135.40	tier - 1	
2331	Id4	1991	Jun	6	1	58571.07	tier - 1	
2332	Id3	1970	?	11	3	60021.40	tier - 1	
2333	Id2	1977	Jun	8	0	62592.87	tier - 2	
2334	Id1	1968	Oct	12	0	63770.43	tier - 1	

	City	tier	State	ID	BMI	HBA1C	Heart	Issues	Any	Transplants
0	tier - 3	R1013	17.580	4.51	No	No	\			
1	tier - 1	R1013	17.600	4.39	No	No				
2	tier - 1	R1013	16.470	6.35	No	No				
3	tier - 3	R1013	17.700	6.28	No	No				
4	tier - 3	R1013	22.340	5.57	No	No				
...
2330	tier - 2	R1012	35.530	5.45	No	No				
2331	tier - 3	R1024	38.095	6.05	No	No				
2332	tier - 1	R1012	34.485	11.87	yes	No				
2333	tier - 3	R1013	30.360	5.77	No	No				
2334	tier - 3	R1013	47.410	7.47	No	No				

	Cancer history	NumberOfMajorSurgeries	smoker	
0	No	1	No	\
1	No	1	No	
2	Yes	1	No	
3	No	1	No	
4	No	1	No	
...	
2330	No	No major surgery	yes	
2331	No	No major surgery	yes	
2332	No	2	yes	
2333	No	No major surgery	yes	
2334	No	No major surgery	yes	

	name
0	German, Mr. Aaron K
1	Rosendahl, Mr. Evan P
2	Albano, Ms. Julie
3	Riveros Gonzalez, Mr. Juan D. Sr.
4	Brietzke, Mr. Jordan
...	...
2330	Kadala, Ms. Kristyn
2331	Osborne, Ms. Kelsey
2332	Lu, Mr. Phil
2333	Lehner, Mr. Matthew D
2334	Hawks, Ms. Kelly

[2335 rows x 17 columns]

```
[98]: df.shape
```

```
[98]: (2335, 17)
```

```
[99]: df.describe()
```

```
[99]:
```

	date	children	charges	BMI	HBA1C
count	2335.000000	2335.000000	2335.000000	2335.000000	2335.000000
mean	15.563597	1.025696	13529.918034	30.972649	6.578998
std	8.720508	1.234754	11898.654299	8.742095	2.228731
min	1.000000	0.000000	563.840000	15.010000	4.000000
25%	8.000000	0.000000	5084.010000	24.600000	4.900000
50%	15.000000	0.000000	9630.910000	30.400000	5.810000
75%	23.000000	2.000000	16912.295000	36.300000	7.955000
max	30.000000	5.000000	63770.430000	55.050000	12.000000

0.2 2. Check for missing values in the dataset

```
[100]: df.isna().sum().sum()
```

```
[100]: 0
```

- There are no missing values in the dataset

0.3 3. Find the percentage of rows that have trivial value (for example, ?), and delete such rows if they do not contain significant information

```
[101]: trivial= df[df == "?"].count(axis=1).sum()
```

```
[102]: trivial
```

```
[102]: 11
```

```
[103]: trivial.shape
```

```
[103]: ()
```

```
[104]: total= df.shape[0]  
total
```

```
[104]: 2335
```

```
[105]: percentage = (trivial / total_rows) * 100
```

```
[106]: percentage
```

```
[106]: 0.47109207708779444
```

- 0.47109 % of rows contain the trivial values.

```
[107]: print("Percentage of trivial rows: {:.2f}%".format(percentage))
```

```
Percentage of trivial rows: 0.47%
```

```
[108]: df = df[df != "?"].dropna()
```

```
[109]: df.shape
```

```
[109]: (2325, 17)
```

0.4 4. Use the necessary transformation methods to deal with the nominal and ordinal categorical variables in the dataset

```
[110]: df_cat = df.select_dtypes(exclude='number')
```

```
[111]: df_cat.columns
```

```
[111]: Index(['Customer ID', 'year', 'month', 'Hospital tier', 'City tier',  
          'State ID', 'Heart Issues', 'Any Transplants', 'Cancer history',  
          'NumberOfMajorSurgeries', 'smoker', 'name'],  
          dtype='object')
```

```
[112]: df['Heart Issues'].value_counts()
```

```
[112]: Heart Issues  
No      1405  
yes      920  
Name: count, dtype: int64
```

```
[113]: df['Any Transplants'].value_counts()
```

```
[113]: Any Transplants  
No      2183  
yes      142  
Name: count, dtype: int64
```

```
[114]: df['Cancer history'].value_counts()
```

```
[114]: Cancer history  
No      1934  
Yes      391  
Name: count, dtype: int64
```

```
[115]: df['smoker'].value_counts()
```

```
[115]: smoker  
No      1839  
yes      486  
Name: count, dtype: int64
```

```
[116]: from sklearn.preprocessing import LabelEncoder  
le= LabelEncoder()
```

```
[117]: df["Heart Issues"]=le.fit_transform(df["Heart Issues"])  
df["Any Transplants"]=le.fit_transform(df["Any Transplants"])  
df["Cancer history"]=le.fit_transform(df["Cancer history"])  
df["smoker"]=le.fit_transform(df["smoker"])
```

```
[118]: df["Heart Issues"].value_counts()
```

```
[118]: Heart Issues
0      1405
1       920
Name: count, dtype: int64
```

- Hospital tier and city tier are ordinal categorical variables

```
[119]: def fun(val):
        return int(val.replace("tier", "").replace(" ", "").replace("-", ""))
```

```
[120]: df['Hospital tier'] = df['Hospital tier'].map(fun)
```

```
[121]: df['City tier'] = df['City tier'].map(fun)
```

```
[122]: df
```

```
[122]:      Customer ID  year month  date  children  charges  Hospital tier
0      Id2335  1992   Jul    9         0    563.84             2 \
1      Id2334  1992  Nov   30         0    570.62             2
2      Id2333  1993   Jun   30         0    600.00             2
3      Id2332  1992   Sep   13         0    604.54             3
4      Id2331  1998   Jul   27         0    637.26             3
...
2329      Id6  1962   Aug    4         0  52590.83             1
2330      Id5  1989   Jun   19         0  55135.40             1
2331      Id4  1991   Jun    6         1  58571.07             1
2333      Id2  1977   Jun    8         0  62592.87             2
2334      Id1  1968  Oct   12         0  63770.43             1

      City tier State ID      BMI  HBA1C  Heart Issues  Any Transplants
0           3    R1013  17.580   4.51         0             0 \
1           1    R1013  17.600   4.39         0             0
2           1    R1013  16.470   6.35         0             0
3           3    R1013  17.700   6.28         0             0
4           3    R1013  22.340   5.57         0             0
...
2329      3    R1011  32.800   6.59         0             0
2330      2    R1012  35.530   5.45         0             0
2331      3    R1024  38.095   6.05         0             0
2333      3    R1013  30.360   5.77         0             0
2334      3    R1013  47.410   7.47         0             0

      Cancer history  NumberOfMajorSurgeries  smoker
0                  0                      1         0 \
1                  0                      1         0
```


2	1	1	0
3	0	1	0
4	0	1	0
...
2329	0	No major surgery	1
2330	0	No major surgery	1
2331	0	No major surgery	1
2333	0	No major surgery	1
2334	0	No major surgery	1

	name
0	German, Mr. Aaron K
1	Rosendahl, Mr. Evan P
2	Albano, Ms. Julie
3	Riveros Gonzalez, Mr. Juan D. Sr.
4	Brietzke, Mr. Jordan
...	...
2329	Baker, Mr. Russell B.
2330	Kadala, Ms. Kristyn
2331	Osborne, Ms. Kelsey
2333	Lehner, Mr. Matthew D
2334	Hawks, Ms. Kelly

[2325 rows x 17 columns]

```
[123]: df['Hospital tier'].value_counts()
```

```
[123]: Hospital tier
2      1334
3       691
1       300
Name: count, dtype: int64
```

0.5 5.The dataset has State ID, which has around 16 states. All states are not represented in equal proportions in the data. Creating dummy variables for all regions may also result in too many insignificant predictors. Nevertheless, only R1011, R1012, and R1013 are worth investigating further. Create a suitable strategy to create dummy variables with these restraints.

```
[124]: df['State ID'].value_counts()
```

```
[124]: State ID
R1013    609
R1011    574
R1012    572
R1024    159
```

```

R1026      84
R1021      70
R1016      64
R1025      40
R1023      38
R1017      36
R1019      26
R1022      14
R1014      13
R1015      11
R1018       9
R1020       6
Name: count, dtype: int64

```

```

[125]: Dummies = pd.get_dummies(df["State ID"], prefix= "State_ID")
Dummies

```

```

[125]:
State_ID_R1011 State_ID_R1012 State_ID_R1013 State_ID_R1014
0             False           False           True           False \
1             False           False           True           False
2             False           False           True           False
3             False           False           True           False
4             False           False           True           False
...
2329          True           False           False           False
2330          False           True           False           False
2331          False           False           False           False
2333          False           False           True           False
2334          False           False           True           False

State_ID_R1015 State_ID_R1016 State_ID_R1017 State_ID_R1018
0             False           False           False           False \
1             False           False           False           False
2             False           False           False           False
3             False           False           False           False
4             False           False           False           False
...
2329          False           False           False           False
2330          False           False           False           False
2331          False           False           False           False
2333          False           False           False           False
2334          False           False           False           False

State_ID_R1019 State_ID_R1020 State_ID_R1021 State_ID_R1022
0             False           False           False           False \
1             False           False           False           False
2             False           False           False           False

```

3	False	False	False	False
4	False	False	False	False
...
2329	False	False	False	False
2330	False	False	False	False
2331	False	False	False	False
2333	False	False	False	False
2334	False	False	False	False

	State_ID_R1023	State_ID_R1024	State_ID_R1025	State_ID_R1026
0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False
...
2329	False	False	False	False
2330	False	False	False	False
2331	False	True	False	False
2333	False	False	False	False
2334	False	False	False	False

[2325 rows x 16 columns]

```
[126]: Dummy = Dummies[['State_ID_R1011', 'State_ID_R1012', 'State_ID_R1013']]
        Dummy
```

```
[126]:      State_ID_R1011  State_ID_R1012  State_ID_R1013
0           False           False           True
1           False           False           True
2           False           False           True
3           False           False           True
4           False           False           True
...
2329         True           False           False
2330         False          True           False
2331         False           False           False
2333         False           False           True
2334         False           False           True
```

[2325 rows x 3 columns]

```
[127]: df = pd.concat([df, Dummy], axis=1)
```

```
[128]: df
```

[128]:

	Customer ID	year	month	date	children	charges	Hospital tier	
0	Id2335	1992	Jul	9	0	563.84	2	\
1	Id2334	1992	Nov	30	0	570.62	2	
2	Id2333	1993	Jun	30	0	600.00	2	
3	Id2332	1992	Sep	13	0	604.54	3	
4	Id2331	1998	Jul	27	0	637.26	3	
...	
2329	Id6	1962	Aug	4	0	52590.83	1	
2330	Id5	1989	Jun	19	0	55135.40	1	
2331	Id4	1991	Jun	6	1	58571.07	1	
2333	Id2	1977	Jun	8	0	62592.87	2	
2334	Id1	1968	Oct	12	0	63770.43	1	

	City tier	State ID	BMI	HBA1C	Heart Issues	Any Transplants	
0	3	R1013	17.580	4.51	0	0	\
1	1	R1013	17.600	4.39	0	0	
2	1	R1013	16.470	6.35	0	0	
3	3	R1013	17.700	6.28	0	0	
4	3	R1013	22.340	5.57	0	0	
...	
2329	3	R1011	32.800	6.59	0	0	
2330	2	R1012	35.530	5.45	0	0	
2331	3	R1024	38.095	6.05	0	0	
2333	3	R1013	30.360	5.77	0	0	
2334	3	R1013	47.410	7.47	0	0	

	Cancer history	NumberOfMajorSurgeries	smoker	
0	0	1	0	\
1	0	1	0	
2	1	1	0	
3	0	1	0	
4	0	1	0	
...	
2329	0	No major surgery	1	
2330	0	No major surgery	1	
2331	0	No major surgery	1	
2333	0	No major surgery	1	
2334	0	No major surgery	1	

	name	State_ID_R1011	State_ID_R1012	
0	German, Mr. Aaron K	False	False	\
1	Rosendahl, Mr. Evan P	False	False	
2	Albano, Ms. Julie	False	False	
3	Riveros Gonzalez, Mr. Juan D. Sr.	False	False	
4	Brietzke, Mr. Jordan	False	False	
...	
2329	Baker, Mr. Russell B.	True	False	

2330	Kadala, Ms. Kristyn	False	True
2331	Osborne, Ms. Kelsey	False	False
2333	Lehner, Mr. Matthew D	False	False
2334	Hawks, Ms. Kelly	False	False

	State_ID_R1013
0	True
1	True
2	True
3	True
4	True
...	...
2329	False
2330	False
2331	False
2333	True
2334	True

[2325 rows x 20 columns]

0.6 6. The variable NumberOfMajorSurgeries also appears to have string values. Apply a suitable method to clean up this variable

```
[129]: df['NumberOfMajorSurgeries'] = df['NumberOfMajorSurgeries'].replace('No major_
↪surgery',0)
```

```
[130]: df['NumberOfMajorSurgeries'].value_counts()
```

```
[130]: NumberOfMajorSurgeries
0      1070
1       961
2       272
3        22
Name: count, dtype: int64
```

0.7 7. Age appears to be a significant factor in this analysis. Calculate the patients' ages based on their dates of birth.

```
[131]: df['year'] = pd.to_datetime(df['year'], format='%Y').dt.year
df['year']
```

```
[131]: 0      1992
1      1992
2      1993
3      1992
4      1998
```

```

...
2329    1962
2330    1989
2331    1991
2333    1977
2334    1968
Name: year, Length: 2325, dtype: int32

```

```
[132]: df['month'] = pd.to_datetime(df['month'], format='%b').dt.month
df['month']
```

```
[132]: 0      7
      1     11
      2      6
      3      9
      4      7
      ..
2329    8
2330    6
2331    6
2333    6
2334   10
Name: month, Length: 2325, dtype: int32

```

```
[133]: df['DateInt'] = df['year'].astype(str) + df['month'].astype(str).zfill(2) +
      df['date'].astype(str).zfill(2)
```

```
[134]: df['DOB'] = pd.to_datetime(df.DateInt, format='%Y%m%d')
```

```
[135]: df.drop(['DateInt'], inplace=True, axis=1)
```

```
[136]: import datetime as dt
current_date = dt.datetime.now()
```

```
[137]: df['Age'] = (((current_date-df.DOB).dt.days)/365).astype(int)
```

```
[138]: df.head()
```

```
[138]: Customer ID  year  month  date  children  charges  Hospital tier  City tier \
0      Id2335  1992     7     9         0    563.84             2           3 \
1      Id2334  1992    11    30         0    570.62             2           1
2      Id2333  1993     6    30         0    600.00             2           1
3      Id2332  1992     9    13         0    604.54             3           3
4      Id2331  1998     7    27         0    637.26             3           3

State ID  BMI  ...  Any Transplants  Cancer history
0      R1013  17.58  ...              0              0 \

```

1	R1013	17.60	...	0	0
2	R1013	16.47	...	0	1
3	R1013	17.70	...	0	0
4	R1013	22.34	...	0	0

	NumberOfMajorSurgeries	smoker	name
0	1	0	German, Mr. Aaron K \
1	1	0	Rosendahl, Mr. Evan P
2	1	0	Albano, Ms. Julie
3	1	0	Riveros Gonzalez, Mr. Juan D. Sr.
4	1	0	Brietzke, Mr. Jordan

	State_ID_R1011	State_ID_R1012	State_ID_R1013	DOB	Age
0	False	False	True	1992-07-09	30
1	False	False	True	1992-11-30	30
2	False	False	True	1993-06-30	29
3	False	False	True	1992-09-13	30
4	False	False	True	1998-07-27	24

[5 rows x 22 columns]

[139]: df.tail()

[139]:

	Customer ID	year	month	date	children	charges	Hospital tier
2329	Id6	1962	8	4	0	52590.83	1 \
2330	Id5	1989	6	19	0	55135.40	1
2331	Id4	1991	6	6	1	58571.07	1
2333	Id2	1977	6	8	0	62592.87	2
2334	Id1	1968	10	12	0	63770.43	1

	City tier	State ID	BMI	...	Any Transplants	Cancer history
2329	3	R1011	32.800	...	0	0 \
2330	2	R1012	35.530	...	0	0
2331	3	R1024	38.095	...	0	0
2333	3	R1013	30.360	...	0	0
2334	3	R1013	47.410	...	0	0

	NumberOfMajorSurgeries	smoker	name	State_ID_R1011
2329	0	1	Baker, Mr. Russell B.	True \
2330	0	1	Kadala, Ms. Kristyn	False
2331	0	1	Osborne, Ms. Kelsey	False
2333	0	1	Lehner, Mr. Matthew D	False
2334	0	1	Hawks, Ms. Kelly	False

	State_ID_R1012	State_ID_R1013	DOB	Age
2329	False	False	1962-08-04	60
2330	True	False	1989-06-19	33

2331	False	False	1991-06-06	31
2333	False	True	1977-06-08	45
2334	False	True	1968-10-12	54

[5 rows x 22 columns]

[140]: df

```
[140]:
```

	Customer ID	year	month	date	children	charges	Hospital tier
0	Id2335	1992	7	9	0	563.84	2 \
1	Id2334	1992	11	30	0	570.62	2
2	Id2333	1993	6	30	0	600.00	2
3	Id2332	1992	9	13	0	604.54	3
4	Id2331	1998	7	27	0	637.26	3
...
2329	Id6	1962	8	4	0	52590.83	1
2330	Id5	1989	6	19	0	55135.40	1
2331	Id4	1991	6	6	1	58571.07	1
2333	Id2	1977	6	8	0	62592.87	2
2334	Id1	1968	10	12	0	63770.43	1

	City tier	State ID	BMI	...	Any Transplants	Cancer history
0	3	R1013	17.580	...	0	0 \
1	1	R1013	17.600	...	0	0
2	1	R1013	16.470	...	0	1
3	3	R1013	17.700	...	0	0
4	3	R1013	22.340	...	0	0
...
2329	3	R1011	32.800	...	0	0
2330	2	R1012	35.530	...	0	0
2331	3	R1024	38.095	...	0	0
2333	3	R1013	30.360	...	0	0
2334	3	R1013	47.410	...	0	0

	NumberOfMajorSurgeries	smoker	name
0	1	0	German, Mr. Aaron K \
1	1	0	Rosendahl, Mr. Evan P
2	1	0	Albano, Ms. Julie
3	1	0	Riveros Gonzalez, Mr. Juan D. Sr.
4	1	0	Brietzke, Mr. Jordan
...
2329	0	1	Baker, Mr. Russell B.
2330	0	1	Kadala, Ms. Kristyn
2331	0	1	Osborne, Ms. Kelsey
2333	0	1	Lehner, Mr. Matthew D
2334	0	1	Hawks, Ms. Kelly

	State_ID_R1011	State_ID_R1012	State_ID_R1013	DOB	Age
0	False	False	True	1992-07-09	30
1	False	False	True	1992-11-30	30
2	False	False	True	1993-06-30	29
3	False	False	True	1992-09-13	30
4	False	False	True	1998-07-27	24
...
2329	True	False	False	1962-08-04	60
2330	False	True	False	1989-06-19	33
2331	False	False	False	1991-06-06	31
2333	False	False	True	1977-06-08	45
2334	False	False	True	1968-10-12	54

[2325 rows x 22 columns]

0.8 8. The gender of the patient may be an important factor in determining the cost of hospitalization. The salutations in a beneficiary's name can be used to determine their gender. Make a new field for the beneficiary's gender.

```
[141]: def gen(x):
        if 'Ms.' in x:
            return 0
        else:
            return 1
```

```
[142]: df['Gender'] = df['name'].map(gen)
```


```
[143]: df['Gender']
```

```
[143]: 0      1
       1      1
       2      0
       3      1
       4      1
       ..
      2329    1
      2330    0
      2331    0
      2333    1
      2334    0
      Name: Gender, Length: 2325, dtype: int64
```

0.9 9. You should also visualize the distribution of costs using a histogram, box and whisker plot, and swarm plot.


```
[144]: # Histogram
sns.histplot(df['charges'])
```

```
[144]: <Axes: xlabel='charges', ylabel='Count'>
```

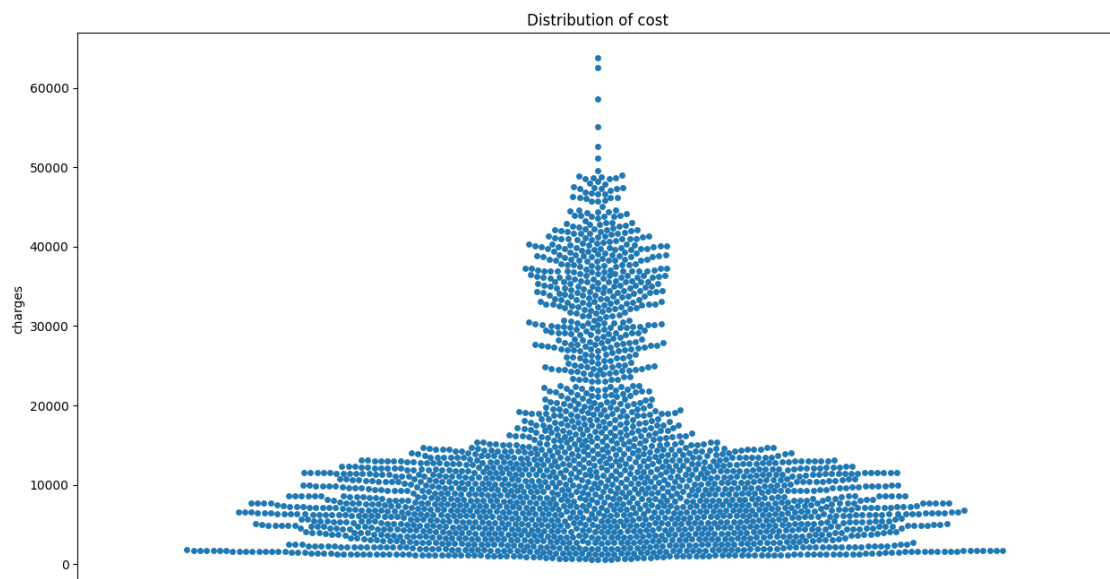


```
[145]: # box and whisker plot
sns.boxplot(df['charges'])
```

```
[145]: <Axes: >
```




```
[148]: # Swarm Plot
plt.figure(figsize=(15,8))
sns.swarmplot(df['charges'])
plt.title('Distribution of cost')
plt.show()
```



0.10 10. State how the distribution is different across gender and tiers of hospitals

```
[147]: sns.countplot(data = df,x = 'Hospital tier', hue = 'Gender')
```

```
[147]: <Axes: xlabel='Hospital tier', ylabel='count'>
```



- In above plot 0 indicates female and 1 indicates male.
- In Hospital tier 1 and 3 we can see that, the count of females are less than compared to the tier 2.

0.11 11. Create a radar chart to showcase the median hospitalization cost for each tier of hospitals

```
[149]: df[df['Hospital tier']==1].charges.median()
```

```
[149]: 32097.434999999998
```

```
[150]: df[df['Hospital tier']==2].charges.median()
```

```
[150]: 7168.76
```

```
[151]: df[df['Hospital tier']==3].charges.median()
```

```
[151]: 10676.83
```

```
[153]: !pip install plotly
```

Collecting plotly

Downloading plotly-5.14.1-py2.py3-none-any.whl (15.3 MB)

```
----- 0.0/15.3 MB ? eta -:--:--
----- 0.2/15.3 MB 11.5 MB/s eta 0:00:02
- ----- 0.4/15.3 MB 4.4 MB/s eta 0:00:04
-- ----- 0.8/15.3 MB 5.7 MB/s eta 0:00:03
--- ----- 1.4/15.3 MB 7.5 MB/s eta 0:00:02
----- 1.9/15.3 MB 8.2 MB/s eta 0:00:02
----- 2.1/15.3 MB 7.4 MB/s eta 0:00:02
----- 2.1/15.3 MB 7.4 MB/s eta 0:00:02
----- 3.0/15.3 MB 6.9 MB/s eta 0:00:02
----- 3.6/15.3 MB 7.4 MB/s eta 0:00:02
----- 4.2/15.3 MB 7.6 MB/s eta 0:00:02
----- 4.2/15.3 MB 7.1 MB/s eta 0:00:02
----- 4.8/15.3 MB 7.2 MB/s eta 0:00:02
----- 5.1/15.3 MB 7.2 MB/s eta 0:00:02
----- 5.7/15.3 MB 7.2 MB/s eta 0:00:02
----- 6.1/15.3 MB 7.5 MB/s eta 0:00:02
----- 6.4/15.3 MB 7.2 MB/s eta 0:00:02
----- 6.5/15.3 MB 6.9 MB/s eta 0:00:02
----- 7.6/15.3 MB 7.6 MB/s eta 0:00:02
----- 7.9/15.3 MB 7.5 MB/s eta 0:00:01
----- 8.5/15.3 MB 7.7 MB/s eta 0:00:01
----- 8.7/15.3 MB 7.7 MB/s eta 0:00:01
----- 9.1/15.3 MB 7.7 MB/s eta 0:00:01
----- 9.5/15.3 MB 7.6 MB/s eta 0:00:01
----- 9.9/15.3 MB 7.5 MB/s eta 0:00:01
----- 10.2/15.3 MB 7.5 MB/s eta 0:00:01
----- 10.7/15.3 MB 7.8 MB/s eta 0:00:01
----- 11.1/15.3 MB 7.8 MB/s eta 0:00:01
----- 11.5/15.3 MB 7.7 MB/s eta 0:00:01
----- 11.8/15.3 MB 7.5 MB/s eta 0:00:01
----- 12.3/15.3 MB 7.4 MB/s eta 0:00:01
----- 12.6/15.3 MB 8.0 MB/s eta 0:00:01
----- 13.0/15.3 MB 7.8 MB/s eta 0:00:01
----- 13.3/15.3 MB 7.7 MB/s eta 0:00:01
----- 13.7/15.3 MB 7.7 MB/s eta 0:00:01
```

```

----- 14.0/15.3 MB 7.5 MB/s eta 0:00:01
----- 14.4/15.3 MB 7.5 MB/s eta 0:00:01
----- 14.7/15.3 MB 7.9 MB/s eta 0:00:01
----- 15.1/15.3 MB 7.6 MB/s eta 0:00:01
----- 15.3/15.3 MB 7.7 MB/s eta 0:00:01
----- 15.3/15.3 MB 7.7 MB/s eta 0:00:01
----- 15.3/15.3 MB 7.7 MB/s eta 0:00:01
----- 15.3/15.3 MB 7.7 MB/s eta 0:00:01
----- 15.3/15.3 MB 7.7 MB/s eta 0:00:01
----- 15.3/15.3 MB 7.7 MB/s eta 0:00:01
----- 15.3/15.3 MB 6.2 MB/s eta 0:00:00
Requirement already satisfied: packaging in c:\users\91805\anaconda1\lib\site-
packages (from plotly) (23.0)
Collecting tenacity>=6.2.0
  Downloading tenacity-8.2.2-py3-none-any.whl (24 kB)
Installing collected packages: tenacity, plotly
Successfully installed plotly-5.14.1 tenacity-8.2.2

```

[155]: `!pip install --upgrade plotly`

```

Requirement already satisfied: plotly in c:\users\91805\anaconda1\lib\site-
packages (5.14.1)
Requirement already satisfied: tenacity>=6.2.0 in
c:\users\91805\anaconda1\lib\site-packages (from plotly) (8.2.2)
Requirement already satisfied: packaging in c:\users\91805\anaconda1\lib\site-
packages (from plotly) (23.0)

```

```

[163]: import plotly.graph_objects as go
import pandas as pd

df1 = df1 = pd.DataFrame(dict(
    r=[32097.434999999998, 7168.76, 10676.83],
    theta=['Tier 1', 'Tier 2', 'Tier 3']
))

df2 = pd.concat([df1, df1.iloc[0]], ignore_index=True)

fig = go.Figure(data=go.Scatterpolar(
    r=df2['r'],
    theta=df2['theta'],
    fill='toself'
))

fig.update_layout(polar=dict(radialaxis=dict(visible=True, range=[0, 35000])))

fig.show()

```

0.12 12. Create a frequency table and a stacked bar chart to visualize the count of people in the different tiers of cities and hospitals

```
[166]: city_freq = df["City tier"].value_counts().rename_axis('City&hospital_tier').  
        ↪reset_index(name='city_counts')
```

```
[167]: hospital_freq = df["Hospital tier"].value_counts().  
        ↪rename_axis('City&hospital_tier').reset_index(name='hospital_counts')
```


```
[168]: freq_table = pd.merge(city_freq, hospital_freq, on = 'City&hospital_tier')
```

```
[169]: freq_table
```

```
[169]:
```

	City&hospital_tier	city_counts	hospital_counts
0	2	807	1334
1	3	789	691
2	1	729	300

```
[170]: x = freq_table['City&hospital_tier']  
        y1 = freq_table['city_counts']  
        y2 = freq_table['hospital_counts']  
  
        # plot bars in stack manner  
        plt.bar(x, y1, color='r')  
        plt.bar(x, y2, bottom=y1, color='b')  
        plt.show()
```



0.13 13. Test the following null hypotheses:

- a. The average hospitalization costs for the three types of hospitals are not significantly different
- b. The average hospitalization costs for the three types of cities are not significantly different
- c. The average hospitalization cost for smokers is not significantly different from the average cost for nonsmokers
- d. Smoking and heart issues are independent

```
[175]: from scipy.stats import ttest_1samp
import scipy.stats as stats
```

a. The average hospitalization costs for the three types of hospitals are not significantly different

```
[176]: print('Null Hypothesis => Average hospitalization costs for the three types of_
↳ hospitals are not significantly different.')
f_val, p_val = stats.f_oneway(df[df['Hospital tier'] == 'tier,1']['charges'],
                              df[df['Hospital tier'] == 'tier,2']['charges'],
                              df[df['Hospital tier'] == 'tier,3']['charges'])
print('P-value :',p_val)
```



```

if p_val < 0.05:
    print("Reject null hypothesis")
else:
    print("Accept null hypothesis")

```

Null Hypothesis => Average hospitalization costs for the three types of hospitals are not significantly different.

P-value : nan

Accept null hypothesis

We already found the median cost of tier 1 hospitals: 32097.434999999998, median cost of tier 2 hospitals: 7168.76 and .median cost of tier 3 hospitals: 10676.83. Interpretation H0: the distributions of all samples are equal and H1: the distributions of one or more samples are not equal

```

[177]: from scipy.stats import friedmanchisquare
data1 = [32097.43]
data2 = [7168.76]
data3 = [10676.83]
stat, p = friedmanchisquare(data1, data2, data3)
print('stat=%.3f, p=%.3f' % (stat, p))
if p > 0.05:
    print('Probably the same distribution')
else:
    print('Probably different distributions')

```

stat=2.000, p=0.368

Probably the same distribution

b. The average hospitalization costs for the three types of cities are not significantly different

```

[178]: print("median cost of tier 1 city:", df[df["City tier"]==1].charges.median())
print("median cost of tier 2 city:", df[df["City tier"]==2].charges.median())
print("median cost of tier 3 city:", df[df["City tier"]==3].charges.median())

```

median cost of tier 1 city: 10027.15

median cost of tier 2 city: 8968.33

median cost of tier 3 city: 9880.07

```

[179]: data1 = [10027.15]
data2 = [8968.33]
data3 = [9880.07]
stat, p = friedmanchisquare(data1, data2, data3)
print('stat=%.3f, p=%.3f' % (stat, p))
if p > 0.05:
    print('Probably the same distribution')
else:
    print('Probably different distributions')

```

stat=2.000, p=0.368
Probably the same distribution

```
[182]: print('Null Hypothesis => Average hospitalization costs for the three types of cities are not significantly different.')
f_val, p_val = stats.f_oneway(df[df['City tier'] == 'tier,1']['charges'],
                             df[df['City tier'] == 'tier,2']['charges'],
                             df[df['City tier'] == 'tier,3']['charges'])

print('P-value :',p_val)
if p_val < 0.05:
    print("Reject null hypothesis")
else:
    print("Accept null hypothesis")
```

Null Hypothesis => Average hospitalization costs for the three types of cities are not significantly different.
P-value : nan
Accept null hypothesis

c. The average hospitalization cost for smokers is not significantly different from the average cost for nonsmokers

```
[180]: print("median cost of smoker:", df[df["smoker"]==1].charges.median())
print("median cost of non smoker:", df[df["smoker"]==0].charges.median())
```

median cost of smoker: 34125.475
median cost of non smoker: 7537.16

```
[186]: from scipy.stats import kruskal
data1 = [34125.475]
data2 = [7537.16]
stat, p = kruskal(data1, data2)
print('stat=%.3f, p=%.3f' % (stat, p))
if p > 0.05:
    print('Probably the same distribution')
else:
    print('Probably different distributions')
```

stat=1.000, p=0.317
Probably the same distribution

```
[184]: print('Null Hypothesis => Average hospitalization costs for smokers is not significantly different from the average cost for nonsmokers.')
t_val, p_val = stats.ttest_ind(df[df['smoker'] == 'yes']['charges'],
                               df[df['smoker'] == 'no']['charges'])

print('P-value :',p_val)
if p_val < 0.05:
```

```

    print("Reject null hypothesis")
else:
    print("Accept null hypothesis")

```

Null Hypothesis => Average hospitalization costs for smokers is not significantly different from the average cost for nonsmokers.

P-value : nan

Accept null hypothesis

- Interpretation H0: the two samples are independent. H1: there is a dependency between the samples.

d. Smoking and heart issues are independent

```

[187]: # d. Smoking and heart issues are independent
from scipy.stats import chi2_contingency
table = [[df["Heart Issues"].value_counts()], [df["smoker"].value_counts()]]
stat, p, dof, expected = chi2_contingency(table)
print('stat=%.3f, p=%.3f' % (stat, p))
if p > 0.05:
    print('Probably independent')
else:
    print('Probably dependent')

```

stat=191.145, p=0.000

Probably dependent

```

[190]: from scipy.stats import chi2_contingency
contingency_table = pd.crosstab(df['smoker'], df['Heart Issues'])
chi2, p, dof, expected = chi2_contingency(contingency_table)
print(f'P-value = {p}')
if p < 0.05:
    print("Reject the null hypothesis, Smoking and heart issues are independent.
    ↪")
else:
    print("Accept null hypothesis, Smoking and heart issues are independent.")

```

P-value = 0.7694797581780767

Accept null hypothesis, Smoking and heart issues are independent.

[]:

1 Machine Learning

1.0.1 1.Examine the correlation between predictors to identify highly correlated predictors. Use a heatmap to visualize this.

```
[191]: df.drop(["Customer ID","State ID",'name', 'year', 'month', 'date', 'DOB'],  
             ↪inplace=True, axis=1)
```

```
[192]: df.head()
```

```
[192]:
```

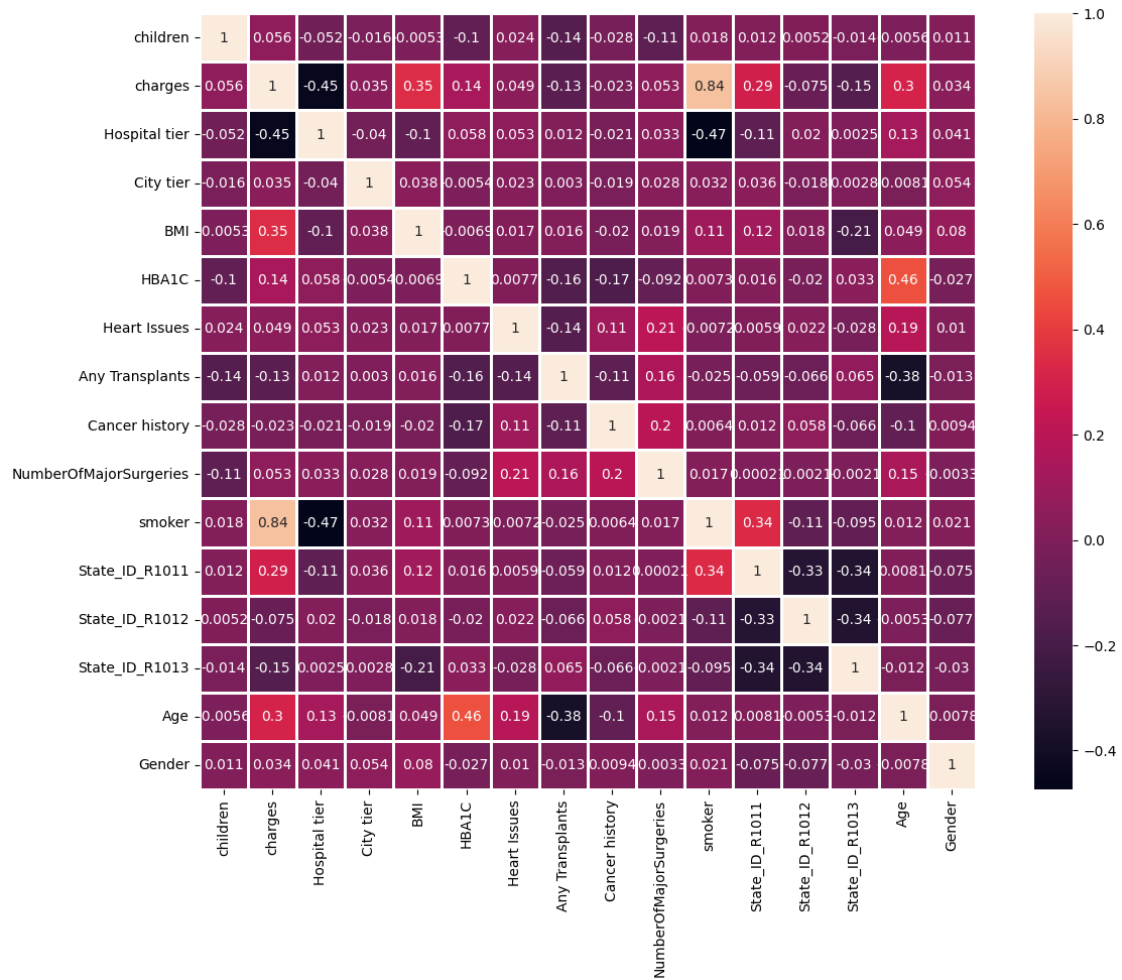
	children	charges	Hospital tier	City tier	BMI	HBA1C	Heart Issues
0	0	563.84	2	3	17.58	4.51	0 \
1	0	570.62	2	1	17.60	4.39	0
2	0	600.00	2	1	16.47	6.35	0
3	0	604.54	3	3	17.70	6.28	0
4	0	637.26	3	3	22.34	5.57	0

	Any Transplants	Cancer history	NumberOfMajorSurgeries	smoker
0	0	0	1	0 \
1	0	0	1	0
2	0	1	1	0
3	0	0	1	0
4	0	0	1	0

	State_ID_R1011	State_ID_R1012	State_ID_R1013	Age	Gender
0	False	False	True	30	1
1	False	False	True	30	1
2	False	False	True	29	0
3	False	False	True	30	1
4	False	False	True	24	1

```
[193]: plt.figure(figsize=(15,10))  
       sns.heatmap(df.corr(),square=True,annot=True,linewidths=1)
```

```
[193]: <Axes: >
```



1.0.2 2. Develop and evaluate the final model using regression with a stochastic gradient descent optimizer. Also, ensure that you apply all the following suggestions:

- Perform the stratified 5-fold cross-validation technique for model building and validation
- Use standardization and hyperparameter tuning effectively
- Use sklearn-pipelines
- Use appropriate regularization techniques to address the bias-variance trade-off
- a. Create five folds in the data, and introduce a variable to identify the folds
- b. For each fold, run a for loop and ensure that 80 percent of the data is used to train the model and the remaining 20 percent is used to validate it in each iteration
- c. Develop five distinct models and five distinct validation scores (root mean squared error values)
- d. Determine the variable importance scores, and identify the redundant variables

```
[194]: from sklearn.model_selection import train_test_split
```

```
[196]: x = df.drop(["charges"], axis=1)
y = df['charges']
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=.
↳20,random_state=10)
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
x_train = sc.fit_transform(x_train)
x_test = sc.fit_transform(x_test)
from sklearn.linear_model import SGDRegressor
```

```
[197]: from sklearn.model_selection import GridSearchCV
params = {'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1, 0.2,0.3,0.4,0.5,
0.6,0.7,0.8,0.9,1.0,2.0,3.0,4.0,5.0,6.0,7.0,8.0,
9.0,10.0,20,50,100,500,1000],
'penalty': ['l2', 'l1', 'elasticnet']}
sgd = SGDRegressor()
# Cross Validation
folds = 5
model_cv = GridSearchCV(estimator = sgd,
param_grid = params,
scoring = 'neg_mean_absolute_error',
cv = folds,
return_train_score = True,
verbose = 1)
model_cv.fit(x_train,y_train)
```

Fitting 5 folds for each of 84 candidates, totalling 420 fits

```
[197]: GridSearchCV(cv=5, estimator=SGDRegressor(),
param_grid={'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1, 0.2, 0.3,
0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0,
4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 20, 50,
100, 500, 1000],
'penalty': ['l2', 'l1', 'elasticnet']},
return_train_score=True, scoring='neg_mean_absolute_error',
verbose=1)
```

```
[198]: model_cv.best_params_
```

```
[198]: {'alpha': 4.0, 'penalty': 'l1'}
```

```
[200]: sgd = SGDRegressor(alpha= 100, penalty= 'l1')
```

```
[201]: sgd.fit(x_train, y_train)
```

```
[201]: SGDRegressor(alpha=100, penalty='l1')
```

```
[202]: sgd.score(x_test, y_test)
```

```
[202]: 0.8711282239533751
```

```
[203]: y_pred = sgd.predict(x_test)
```

```
[205]: from sklearn.metrics import mean_squared_error, mean_absolute_error
sgd_mae = mean_absolute_error(y_test, y_pred)
sgd_mse = mean_squared_error(y_test, y_pred)
sgd_rmse = sgd_mse*(1/2.0)
```

```
[206]: print("MAE:", sgd_mae)
print("MSE:", sgd_mse)
print("RMSE:", sgd_rmse)
```

```
MAE: 2681.183902514782
MSE: 18756962.940102503
RMSE: 9378481.470051251
```

```
[207]: importance = sgd.coef_
pd.DataFrame(importance, index = x.columns, columns=['Feature_imp'])
```

```
[207]:
```

	Feature_imp
children	324.625902
Hospital tier	-1135.536560
City tier	0.000000
BMI	2697.227677
HBA1C	138.011477
Heart Issues	0.000000
Any Transplants	0.000000
Cancer history	0.000000
NumberOfMajorSurgeries	0.000000
smoker	8998.720349
State_ID_R1011	0.000000
State_ID_R1012	0.000000
State_ID_R1013	-218.159993
Age	3401.726864
Gender	0.000000

1.1 3. Use random forest and extreme gradient boosting for cost prediction, share your cross validation results, and calculate the variable importance scores

```
[208]: from sklearn.ensemble import RandomForestRegressor
```

```
[209]: rf = RandomForestRegressor(n_estimators = 1000, random_state = 42)
```

```
[210]: rf.fit(x_train, y_train)
```

```
[210]: RandomForestRegressor(n_estimators=1000, random_state=42)
```

```
[211]: score = rf.score(x_test,y_test)
score
```

```
[211]: 0.901740825808739
```

```
[212]: y_pred = rf.predict(x_test)
```

```
[213]: rf_mae = mean_absolute_error(y_test, y_pred)
```

```
[214]: rf_mae
```

```
[214]: 1940.0160889462331
```

Extreme gradient boosting

```
[215]: from sklearn.ensemble import GradientBoostingRegressor
```

```
[216]: gbr = GradientBoostingRegressor(n_estimators = 1000, random_state = 42)
```

```
[217]: gbr.fit(x_train, y_train)
```

```
[217]: GradientBoostingRegressor(n_estimators=1000, random_state=42)
```

```
[218]: score = gbr.score(x_test,y_test)
score
```

```
[218]: 0.8650084426981199
```

```
[219]: y_pred = gbr.predict(x_test)
```

```
[220]: gbr_mae = mean_absolute_error(y_test, y_pred)
gbr_mae
```

```
[220]: 2424.08578378777
```

- Since Mean Absolute Error of Random Forest is less than extreme gradient boosting, Random Forest algorithm works well.

- 1.1.1 4. Case scenario: Estimate the cost of hospitalization for Christopher, Ms. Jayna (her date of birth is 12/28/1988, height is 170 cm, and weight is 85 kgs). She lives in a tier-1 city and her state's State ID is R1011. She lives with her partner and two children. She was found to be nondiabetic ($HbA1c = 5.8$). She smokes but is otherwise healthy. She has had no transplants or major surgeries. Her father died of lung cancer. Hospitalization costs will be estimated using tier-1 hospitals.

```
[222]: df.columns
```

```
[222]: Index(['children', 'charges', 'Hospital tier', 'City tier', 'BMI', 'HBA1C',  
        'Heart Issues', 'Any Transplants', 'Cancer history',  
        'NumberOfMajorSurgeries', 'smoker', 'State_ID_R1011', 'State_ID_R1012',  
        'State_ID_R1013', 'Age', 'Gender'],  
        dtype='object')
```

```
[223]: date = str(19881228)  
date1 = pd.to_datetime(date, format = "%Y%m%d")
```

```
[226]: current_date = dt.datetime.now()  
current_date
```

```
[226]: datetime.datetime(2023, 5, 25, 18, 55, 19, 779951)
```

```
[227]: age = (current_date - date1)  
age
```

```
[227]: Timedelta('12566 days 18:55:19.779951')
```

```
[228]: age = int(12566/365)  
age
```

```
[228]: 34
```

- So the age of Christopher, Ms. Jayna is 34

```
[229]: height_m = 170/100  
height_sq = height_m*height_m  
BMI = 85/height_sq  
np.round(BMI,2)
```

```
[229]: 29.41
```

- BMI is 29.41

```
[268]: df.columns
```

```
[268]: Index(['children', 'Hospital tier', 'City tier', 'BMI', 'HBA1C',
          'Heart Issues', 'Any Transplants', 'Cancer history',
          'NumberOfMajorSurgeries', 'smoker', 'State_ID_R1011', 'State_ID_R1012',
          'State_ID_R1013', 'age', 'gender'],
          dtype='object')
```

```
[270]: list = [[2,1,1,29.41,5.8,0,0,0,0,1,1,0,0,34,0]]
```

```
[271]: df = pd.DataFrame(list, columns = ['children','Hospital tier', 'City tier', 'BMI', 'HBA1C', 'Heart Issues', 'Any Transplants',
          'Cancer history', 'NumberOfMajorSurgeries', 'smoker', 'State_ID_R1011', 'State_ID_R1012',
          'State_ID_R1013', 'age', 'gender'] )
df
```

```
[271]:
```

	children	Hospital tier	City tier	BMI	HBA1C	Heart Issues	
0	2	1	1	29.41	5.8	0	\
	Any Transplants	Cancer history	NumberOfMajorSurgeries	smoker			
0	0	0	0	0	1		\
	State_ID_R1011	State_ID_R1012	State_ID_R1013	age	gender		
0	1	0	0	34	0		

5. Find the predicted hospitalization cost using all models. The predicted value should be the mean of the five models' predicted values.

```
[274]: Hospital_cost = []
```

```
[275]: Cost1 = sgd.predict(df)
Hospital_cost.append(Cost1)
Cost2 = rf.predict(df)
Hospital_cost.append(Cost2)
Cost3 = gbr.predict(df)
Hospital_cost.append(Cost3)
avg_cost = np.mean(Hospital_cost)
avg_cost
```

```
[275]: 104554.8763667323
```

- The mean of the five models' predicted average value is 104554.8763667323

```
[ ]:
```